Parallel Speculative Decoding with Adaptive Draft Length

Tianyu Liu[1,3*] **Yun Li**[2†] **Qitan Lv**[1] **Kai Liu**[2] **Jianchen Zhu**[2] **Winston Hu**[2]

[1]University of Science and Technology of China
[2]Tencent
[3]OpenGVLab, Shanghai AI Laboratory
{tianyu_liu, qitanlv}@mail.ustc.edu.cn
yunli.charles@gmail.com
{raccoonliu, dickzhu, winstony}@tencent.com

Abstract

Speculative decoding (SD), where an extra draft model is employed to provide multiple *draft* tokens first and then the original target model verifies these tokens in parallel, has shown great power for LLM inference acceleration. However, existing SD methods suffer from the mutual waiting problem, i.e., the target model gets stuck when the draft model is *guessing* tokens, and vice versa. This problem is directly incurred by the asynchronous execution of the draft model and the target model, and is exacerbated due to the fixed draft length in speculative decoding. To address these challenges, we propose a conceptually simple, flexible, and general framework to boost speculative decoding, namely **P**arallel sp**E**culative decoding with **A**daptive d**R**aft **L**ength (PEARL). Specifically, PEARL proposes *pre-verify* to verify the first draft token in advance during the drafting phase, and *post-verify* to generate more draft tokens during the verification phase. PEARL parallels the drafting phase and the verification phase via applying the two strategies, and achieves adaptive draft length for different scenarios, which effectively alleviates the mutual waiting problem. Moreover, we theoretically demonstrate that the mean accepted tokens of PEARL is more than existing *draft-then-verify* works. Experiments on various text generation benchmarks demonstrate the effectiveness of our PEARL, leading to a superior speedup performance up to **3.79**× and **1.52**×, compared to auto-regressive decoding and vanilla speculative decoding, respectively.

1 Introduction

Large language models (LLMs) such as GPT-4, LlaMA, and Claude 3 [Achiam et al., 2023, cla, Bommasani et al., 2021, Touvron et al., 2023] have dominated the natural language understanding and generation [Khurana et al., 2023] over a wide range of applications. However, the substantial inference latency of these LLMs has emerged as a significant obstacle bounding their broader application in scenarios with restricted computational resources. This latency primarily originates from the auto-regressive token-by-token decoding process wherein decoding K tokens requires K serial runs of LLMs, incurring exacerbated latency with both the length of generated tokens and the model scale.

To address this challenge, extensive research efforts have been devoted to accelerating LLM inference. One innovative inference paradigm, Speculative Decoding (SD), has emerged as a new trend and shown superior performance. As shown in Figure. 1(a), the key idea of SD algorithm is to employ an
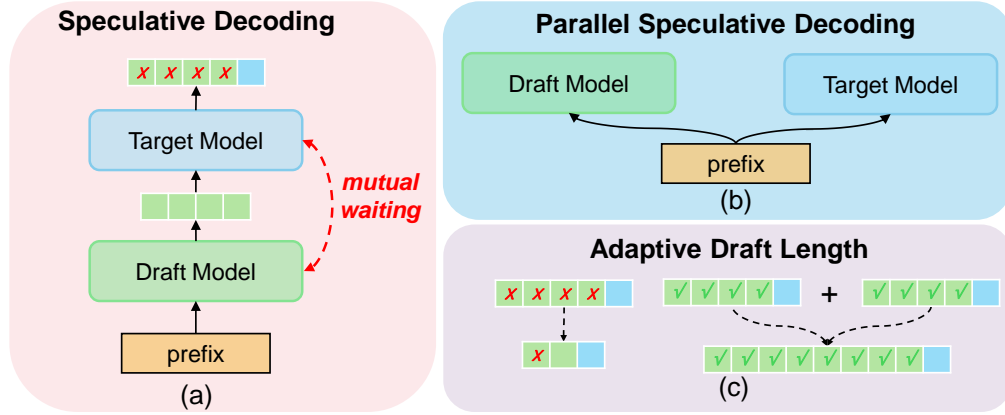

Figure 1: An overview of speculative decoding and our PEARL. (a) SD employs a draft model to provide multiple drafts and then the target model verifies the drafts in parallel. However, SD suffers from the mutual waiting problem, i.e., the target model gets stuck when the draft model is *guessing* tokens, and vice versa. (b) PEARL parallels the drafting and verification process to alleviate the mutual waiting problem. (c) PEARL can utilize adaptive draft length to further mitigate the mutual waiting problem. PEARL can generate less draft tokens if they will be rejected (left in (c)), while PEARL can generate more draft tokens if they can be accepted (right in (c)).

extra small model (referred as the *draft model*) to firstly generate γ draft tokens for the original large model (referred as the *target model*), and then the target model verifies these draft tokens in parallel within a single forward. Here, γ is a fixed hyperparameter **window size**. **Draft length** is the number of tokens generated by the draft model in a continuous execution. Therefore, the draft length is set to γ in SD. Following-up works effectively extend this framework by either removing the necessity of the draft model [Cai et al., 2024, Fu et al., 2024, Zhang et al., 2023] or identifying compact draft model with high distribution alignment [Zhou et al., 2023, Zhao et al., 2024, Miao et al., 2023]. Extensive experiments demonstrate that this *draft-then-verify* framework effectively enhances the concurrency of the target model, thereby significantly accelerating the inference process.

Albeit with multiple benefits of this *draft-then-verify* framework, it confronts one significant challenge that may hinder its performance and deployment—the mutual waiting problem. That is, the target model will be idle when the draft model is generating the draft tokens and the draft model will be idle when the target model is verifying the previously drafted tokens. This mutual waiting problem primarily stems from two limitations inherent in speculative decoding: **(i)** the asynchronous execution of the draft and verify phases, which directly results in the mutual waiting problem; and **(ii)** the fixed draft length, which cannot adapt to most decoding steps and thus exacerbate the mutual waiting problem.

Therefore, in this paper, we seek to answer the question: *Can we draft and verify in parallel and adaptively adjust draft length?* With this consideration, we propose a conceptually simple, flexible, and general framework to boost speculative decoding, namely **P**arallel sp**E**culative decoding with **A**daptive d**R**aft **L**ength (PEARL). Specifically, PEARL consists of two strategies *pre-verify* and *post-verify*: **(i)** *pre-verify* uses the target model to verify the first draft token during drafting phase, which allows the draft model to generate less draft tokens in difficult scenarios; **(ii)** *post-verify* uses the draft model to continue generating draft tokens during verification phase, which provides more draft tokens in simple situations. As shown in Figure.1(b)(c), PEARL effectively alleviates the mutual waiting problem with **parallelism** and **adaptive draft length** via these two strategies. Moreover, we theoretically show that PEARL can **eliminate the burden of tuning the hyperparameter** γ, and the mean accepted tokens of PEARL is **more than** existing *draft-then-verify* works.

Our key contributions can be summarized as follows:

(i) We propose PEARL, a novel inference acceleration framework, which can effectively alleviate the mutual waiting problem with parallelism and adaptive draft length.

Figure 2: Motivated experiments. (a) The time of both drafting phase and verification phase is non-negligible, therefore the asynchronous execution of the draft model and the target model directly incurs the mutual waiting problem. (b) We observe that the optimal draft length changes significantly in different iterations, which exacerbates the mutual waiting problem.

(ii) We theoretically derive the optimal window size and the mean accepted tokens of our PEARL, which demonstrates the effectiveness of our PEARL.

(iii) We conduct extensive experiments on various text generation benchmarks, leading to a superior speedup performance up to **3.79**× and **1.52**×, compared to auto-regressive decoding and speculative decoding, respectively.

2 Background

Notations. In this paper, we use M_q to denote the draft model and M_p to denote the target model. $M_q(\cdot), M_p(\cdot)$ denotes the logits of the next token of a single forward of M_q, M_p respectively. γ is a hyperparameter to control the window size during speculative decoding. We denote the running speed between M_q and M_p as c, which is defined as the ratio between the time for a single forward of M_p and the time for a single forward of M_q, i.e., $c = T(M_p(\cdot))/T(M_q(\cdot))$.

Speculative decoding. Given an input sequence $\mathbf{x}$ as prefix, a speculative decoding step consists of a drafting phase and a verification phase. During the drafting phase, the draft model M_q is employed to give γ draft tokens $x_1, x_2, ..., x_\gamma$ by running γ times model forward and sample. Here, we denote $M_q(\mathbf{x} + [x_1, ..., x_{i-1}])$ as q_i, then each draft token is given by $x_i \sim q_i, i = 1, ..., \gamma$. During the verification phase, the prefix $\mathbf{x}$ together with γ draft tokens are sent to M_p for verification. The target model M_p inputs $\mathbf{x} + [x_1, ..., x_\gamma]$ and outputs the logits $p_1, p_2, ..., p_{\gamma+1}$. Then SD sequentially verifies x_i via speculative sampling, where the acceptance rate is given by:

$$\alpha_i = \begin{cases} 1 & p_i[x_i] \geq q_i[x_i], \\ \dfrac{p_i[x_i]}{q_i[x_i]} & p_i[x_i] < q_i[x_i], \end{cases} \tag{1}$$

If SD rejects x_i, it will resample a token from $norm(\max(0, p_i - q_i))$, otherwise SD accepts all the draft tokens and samples an additional token from $p_{\gamma+1}$. In this way, each SD step generates tokens with a number of at least 1 and at most $\gamma + 1$, leading to the efficiency acceleration.

Window size and draft length. We emphasize that the window size is a hyperparameter that controls the drafting behavior. Draft length is the number of tokens generated by the draft model in a continuous execution, which is fixed and same as the window size in SD, while draft length is adaptive and may be not equal to window size in PEARL.

Choices of model combinations. There are several ways to obtain the draft model for speculative inference. On the one hand, existing open-source LLMs typically include corresponding off-the-shelf smaller architectures pre-trained on the same datasets (e.g., Codellama 7B as the draft model of

3

Codellama 34B [Roziere et al., 2023]). On the other hand, it is feasible to train a highly aligned draft model. However, the training process is typically resource-consuming, and the resulting model often lacks generalizability across different tasks. Therefore, in this paper, we focus primarily on leveraging existing open-source model families to assess our PEARL.

3 Methodology

3.1 Motivated Observation

We conduct some experiments to directly illustrate the mutual waiting problem. As shown in Figure. 2(a), we observe that the time consumed during the drafting phase and the verification phase is usually non-negligible. Take the instance of Codellama 7B & 70B, at each decoding step, the draft model will be inactive for 0.07s during the verification phase, while the target model will be inactive for 0.14s during the drafting phase. Therefore, the mutual waiting problem exists **at any timestamp**.

The asynchronous execution of the draft model and the target model is the direct cause of the mutual waiting problem, which is determined by two requirements of speculative decoding: (1) the drafting phase requires the input prefix to be verified; (2) the verification phase requires the draft model to complete generating draft tokens. This implies the great potential for alleviating the mutual waiting problem through parallelism: if we can remove the two requirements and parallel the drafting phase and the verification phase, a substantial acceleration can be possible.

Another limitation that aggravates the mutual waiting problem is the fixed draft length in SD, which is not appropriate for all the decoding steps. As shown in Figure 2(b), the optimal draft length changes significantly in different iterations. On the one hand, when the optimal draft length is less than the fixed draft length, the draft model will generate meaningless draft tokens that stucks the target model. On another hand, when the optimal draft length is more than the fixed draft length, the draft model could have generated more draft tokens that can be accepted by the target model with a single forward. However, a fixed draft length will interrupt the longer drafting phase and take an additional verification phase, which strengths the mutual waiting problem as well. This motivates our PSD to further alleviates the mutual waiting problem with adaptive draft length.

Together with the two motivations, we propose two simple and effective strategies, *pre-verify* and *post-verify*. The *pre-verify* removes requirement 2 and allows the target model to verify the first draft token in advance. The *post-verify* removes requirement 1 and allows the draft model to continue generating draft tokens during the verification phase. The two strategies enable parallelism and achieve adaptive draft length to effectively alleviate the mutual waiting problem.

3.2 Pre-verify: verify the first draft token in advance.

The *pre-verify* strategy aims at removing the requirement that the verification phase requires the draft model to complete generating draft tokens. Therefore, we seek to verify some draft tokens in advance during drafting phase. We delve explicitly into the drafting stage. During the drafting phase, the draft model tries to give γ draft tokens by running γ times model forward. We find that the input of the draft model in γ times forward is $\mathbf{x}$, $\mathbf{x} + [x_1]$, ..., $\mathbf{x} + [x_1, x_2, ..., x_{\gamma-1}]$, respectively. Only the origin prefix $\mathbf{x}$ can be acquired by the target model for parallel verification. Therefore, we propose to run the target model to output the logits $M_p(\mathbf{x})$ in parallel. In this way, we can verify the first token x_1 before the verification phase. We implement the same lossless verification method following [Leviathan et al., 2023] as illustrated in Section 2.

By applying such a *pre-verify* strategy, we can verify the first draft token before the verification phase. If the first token is rejected, all of the following draft tokens are meaningless and should be dropped. Hence we could skip the verification phase and directly conduct the next drafting phase with the prefix $\mathbf{x} + [y_1]$. If the first token is accepted, all the draft tokens will be sent to the target model in the verification phase. In Figure. 3, at the timestamp of $T = 0$, the draft model generates x_1, x_2, x_3 while the target model outputs p_0^t, rejects the first token x_1 and sample another token y_1. At the timestamp of $T = 3t$, the draft model generates x_4, x_5, x_6 while the target model accepts the first token x_4. Then x_4, x_5, x_6 is sent to the target model in the next verification phase.
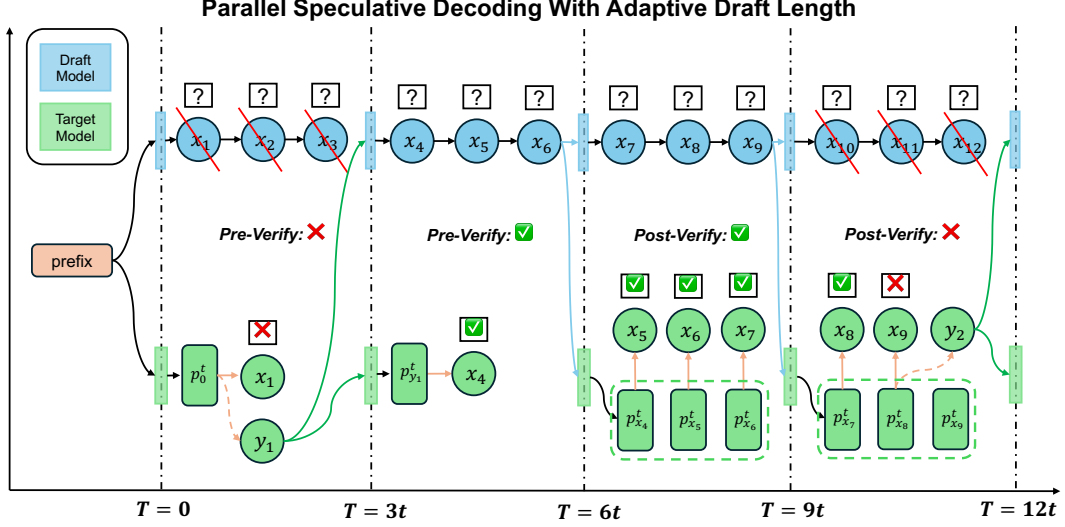
Figure 3: Illustration of our PEARL. At $T = 0$, M_q generates x_1, x_2, x_3 and M_p rejects x_1 with the *pre-verify* strategy. At $T = 3t$, M_p accepts x_4 and switches to the *post-verify* strategy. At $T = 6t$, M_p accepts all draft tokens x_4, x_5, x_6 in the last decoding step, while M_q continues drafting x_7, x_8, x_9. At $T = 9t$, M_p rejects x_9, drops x_{10}, x_{11}, x_{12} and switches to the *pre-verify* strategy. The final output is $[y_1, x_4, x_5, x_6, x_7, x_8, y_2]$.

3.3   Post-verify: continue drafting during verification.

The *post-verify* strategy aims at removing the requirement that the drafting phase requires the input prefix to be verified. However, this assumption brings the limitation that the draft model should be stuck until the target model finishes verification.

Therefore, we discard this assumption and make another assumption: we directly assume that all the draft tokens can be accepted. In this way, We find that when all the γ draft tokens are accepted, sampling a new token from $M_p(\mathbf{x} + [x_1, ..., x_\gamma])$ is not necessary, as the draft model could have generated more draft tokens that can be accepted. Hence we can use the draft model to continue drafting $x_{\gamma+1}, ..., x_{2\gamma}$ during the verification phase.

If all the γ draft tokens are accepted, we can skip the next drafting phase as we have already get the draft tokens in the next drafting phase. The last logit $M_p(\mathbf{x} + [x_1, ..., x_\gamma])$ can be used to verify $x_{\gamma+1}$, which is a *"pre-verify"* process as well. In Figure. 3, at the timestamp of $T = 6t$, the target model takes in x_4, x_5, x_6 and outputs $p_{x_4}^t, p_{x_5}^t, p_{x_6}^t$, while the draft model continues to guess next draft tokens x_7, x_8, x_9. Fortunately, all the draft tokens are accepted, and we can directly conduct next verification phase with prefix $\mathbf{x} + [y_1, x_4, x_5, x_6, x_7, x_8, x_9]$. At the timestamp of $T = 9t$, the target model takes in x_7, x_8, x_9 and outputs $p_{x_7}^t, p_{x_8}^t, p_{x_9}^t$, while the draft model continues to guess the next draft tokens x_{10}, x_{11}, x_{12}. Unfortunately, only x_8 is accepted, and the draft tokens x_{10}, x_{11}, x_{12} will be dropped. Finally, the prefix $\mathbf{x} + [y_1, x_4, x_5, x_6, x_7, x_8, y_2]$ is input to the next drafting phase.

3.4   PEARL: parallel speculative decoding with adaptive draft length

Take together the two strategies, our PEARL framework consists of a draft model, a target model and two strategies to decode tokens. The two strategies are switched according to the verification results in the last decoding step. Algorithm. 1 provides a summary of our PEARL. We also provide more details in Algorithm. 2. Then we show how our PEARL achieves parallelism and adaptive draft length to alleviate the mutual waiting problem.

Parallelism. With the two strategy *pre-verify* and *post-verify*, At any timestamp, the draft model and the target model are running in parallel, which directly breaks the asynchronous execution of the draft model and the target model.

5

Adaptive draft length. In our PEARL, the drafting process can be seen as segmented drafting process. If the draft model cannot generate any "right" tokens, the *pre-verify* strategy will avoid the additional drafting process. If the draft model could have generated more "right" tokens, the target model will not interrupt the drafting phase, where the draft model can generate more draft tokens with *post-verify* strategy. Therefore, PEARL can utilize the two simple yet effective strategies to implement adaptive draft length to alleviate the mutual waiting problem.

4 Analysis

In this section, we give some interesting theoretical findings to demonstrate the generalization ability and effectiveness of our PEARL.

4.1 Eliminating the burden of tuning γ

During a standard speculative decoding step, the draft model takes γ model forward to decode γ draft tokens and then the target model takes 1 model forward to verify the draft tokens. If γ is set too small, we cannot fully exploit the ability of the draft model, and the realistic speedup will be tiny. If γ is set too large, the extra overhead of running draft model will cover the acceleration of speculative decoding, leading to an undesirable speedup. Hence it is crucial to find an optimal value of γ for considerable acceleration.

Algorithm 1 Parallel Speculative Decoding with Adaptive Draft Length.

Require: the draft model M_q, the target model M_p, the input prefix $\mathbf{x}$, the max generate tokens L, the window size γ.
  Initialization: mode $\leftarrow$ "pre-verify"
  while $len(\mathbf{x}) < L$ **do**
    if mode = "pre-verify" **then**
      $\mathbf{x}$, mode $\leftarrow$ Pre-verify($M_q, M_p, \mathbf{x}, \gamma$)
    else
      $\mathbf{x}$, mode $\leftarrow$ Post-verify($M_q, M_p, \mathbf{x}, \gamma$)
    end if
  end while

Assuming the acceptance rate is α, the optimal value γ' of γ in SD is given as follows:

$$\gamma' = \arg\max_{\gamma} SD(\gamma) = \arg\max_{\gamma} \frac{1 - \alpha^{\gamma+1}}{(1-\alpha)(\frac{\gamma}{c} + 1)} \tag{2}$$

However, there lies some issues in finding γ' with this equation. On the one hand, this equation makes an important assumption that the acceptance rate of each token is i.i.d.. Actually, this assumption is only an approximation, leading to difficulties in finding an appropriate γ. On another hand, the value of α may vary in different scenarios, as the draft model and the target model may show different alignments in different scenarios. Therefore, for a specific downstream application, one needs to search an optimal value of γ, which brings severe burden and limits the application of speculative decoding.

In our PEARL, γ' can be theoretically found without these issues. We give Theorem. 1 to show that the optimal value of γ is exactly c.

Theorem 1 *Given a draft model M_q and a target model M_p, the optimal value of the window size γ is the ratio of the running speed of the draft model and the target model, i.e.,*

$$\gamma' = \arg\max_{\gamma} \ PEARL(\gamma) = c. \tag{3}$$

We prove this theorem in Appendix B. Intuitively, increasing or decreasing γ' does not contribute for better efficiency. We conduct experiments in Section 5.4.1 to empirically demonstrate this theorem, which is important for our PEARL to **eliminate the burden of tuning γ.**

4.2 Expectation of the number of accepted tokens

As our PEARL is based on the standard speculative decoding, we theoretically compare the decoding step of our PEARL and standard SD. Suppose the draft model M_q and the target model M_p are same in the two algorithms. We do not introduce any extra training or fine-tuning, hence the acceptance rate for each token α is same and can be seen as a constant in the following analysis.

We begin our analysis by computing the expectation of the number of accepted tokens within a drafting phase and a verification phase. Note that in our PEARL, if all the draft tokens in a drafting phase are accepted, PEARL will continue drafting more draft tokens. We regard the process that the draft model decodes draft tokens until some tokens are rejected as a drafting phase when computing the expectation.

Theorem 2 *Assuming the acceptance rate of each draft token is α, and α is i.i.d., the expectation of the number of accepted tokens of PEARL is*

$$E(\#accepted\ tokens) = \frac{1}{1-\alpha} + 1. \tag{4}$$

We prove this theorem in Appendix B. We give a more intuitive explanation of this theorem: given a prefix, if the draft model M_q can decode δ draft tokens at most that can be accepted by the target model, then M_q will continue to decode δ tokens without been interrupted by the target model verification. It can be explained as an adaptive "γ" for each prefix to achieve better speedup.

Note that the expectation of accepted tokens of other *draft-then-verify* works is $\frac{1-\alpha^{\gamma+1}}{1-\alpha}$. With Theorem. 2, it is easily to show that **the expectation of accepted tokens of PEARL is more than standard SD**, i.e., $\frac{1}{1-\alpha} + 1 \geq \frac{1-\alpha^{\gamma+1}}{1-\alpha}$. From this perspective, we can demonstrate the effectiveness of our PEARL. We will conduct experiments in Section 5.4.2 to show the empirical results of mean accepted tokens in both standard SD and our PEARL.

5 Experiments

5.1 Experimental setup

Tasks and Datasets. We conduct experiments on various text generation tasks to evaluate the effectiveness of our PEARL, including code generation, arithmetic reasoning, and multi-round dialogue. For code generation task, we utilize HumanEval [Chen et al., 2021], which is composed of 164 code generation problems. For arithmetic reasoning, we employ the GSM8K [Cobbe et al., 2021] dataset with the first 100 math problems. For multi-round dialogue, we use MT-bench [Zheng et al., 2024] with 80 multi-round conversation problems in different scenarios.

Models and Metrics. We evaluate the effectiveness of our PEARL with some state-of-the-art LLMs. For HumanEval, we employ both code-specialized LLMs and general LLMs, including CodeLlama-hf [Roziere et al., 2023], Deepseek-Coder[Guo et al., 2024] and Llama 2[Touvron et al., 2023]. For GSM8K and MT-bench, we employ Llama 2 to conduct text generation. In our experiments, the models with size of 1.3B and 6/7B are used as the draft models and the models with size of 33/34B and 70B are used as the target models. We report the text generation speed to compute the walltime speedup ratio as the metric. We provide more evaluation details in Appendix C.

Baseline Methods. We implement *four* variants of LLM inference methods as our baselines. **(i) Speculate decoding:** standalone SD methods [Leviathan et al., 2023, Chen et al., 2023] resort to a draft model to draft future tokens and then verify them in parallel. **(ii) Ouroboros:** ouroboros [Zhao et al., 2024] proposes phrase candidate pool from the verification process to generate more precise and longer drafts. **(iii) Lookahead Decoding:** look ahead decoding [Fu et al., 2024] caches the generation trajectory (n-grams) as drafts to reduce the number of total decoding steps. **(iv) Distillspec:** DistillSpec [Zhou et al., 2023] distillates and derives a better aligned and compact draft model.

5.2 Main results.

Table 1: Experiment results on the code generation task. The results of speculative decoding and Ouroboros are taken from Zhao et al. [2024]. We **bold** the best results for each model combination. * Some results of ouroboros and lookahead decoding are reproduced in our implementation.

	Codellama7B&34B	Codellama7B&70B	Llama 2 7B&70B	Deepseek 7B&33B	Deepseek 1.3B&33B
Auto Regressive	1.00×	1.00×	1.00×	1.00×	1.00×
Speculative Decoding	1.69×	2.76×	2.14×	2.07×	2.54×
Ouroboros	2.14×	* 3.28×	* 2.10×	* 2.66×	* 3.25×
Lookahead	1.72×	* 1.04×	* 1.01×	* 1.05×	* 1.05×
Ours	**2.35×**	**3.79×**	**3.01×**	**2.75×**	**3.48×**

7

Table 2: Experiment results using Llama 7B&70B on GSM8K and MT-bench. The results of speculative decoding and Ouroboros are taken from Zhao et al. [2024]. The symbol ' - ' denotes that the results has not been report in original papers. We **bold** the best results.

	Speculative Decoding	Ouroboros	Lookahead Decoding	Distillspec	PEARL (**ours**)
GSM8K	1.99×	2.68×	1.00×	-	**2.87×**
MT-bench	1.63×	1.40×	1.45×	2.13×	**2.48×**

Table 1 shows that PEARL significantly outperforms vanilla speculative decoding, ouroboros, and lookahead decoding in all backbone model configurations on the HumanEval dataset. Specifically, PEARL can achieve up to 3.79 × speed up compared with vanilla auto-regressive methods. These experimental results demonstrate that PEARL effectively addresses the mutual waiting issue, thereby achieving significant inference acceleration results compared to methods based on the traditional draft-then-verify framework. PEARL can also achieve significant inference acceleration in conventional natural language processing tasks, whereas vanilla speculative decoding, lookahead decoding yield only limited improvements acceleration. As shown in Table 2, PEARL leads to superior performances on the GSM8K and MT-bench datasets.

5.3 Ablation studies

Table 3: Ablation results of PEARL on HumanEval and GSM8K datasets on the speed up metric.

	HumanEval			GSM8K
Methods	CodeLlama 7B&34B	CodeLlama 7B&70B	DeepSeek 1.3B&33B	Llama 2 7B&70B
PEARL *w/o pre-verify*	2.21×	3.53×	3.19×	2.51×
PEARL *w/o post-verify*	1.64×	2.57×	2.37×	2.15×
PEARL	**2.35×**	**3.79×**	**3.48×**	**2.87×**

To provide more insights of the two proposed strategies, we conduct the ablation study. We denote PEARL without *pre-verify* as PEARL *w/o pre-verify* and PEARL without *post-verify* as PEARL *w/o pre-verify* and present the main results of ablation studies.

As shown in Table 3, the absence of any strategy of PEARL results in a performance degradation of the entire framework. The absence of the *post-verify* strategy exhibits a pronounced impact on the performance of PEARL than the *pre-verify* strategy. Intuitively, the *pre-verify* strategy makes more contributions when the acceptance rate is relatively low, while the *post-verify* strategy makes more contributions when the two model is effectively aligned. Therefore, the two strategies are complementary and accelerate inference together. In our experiments, all the model combinations show great alignment, which leads to the superiority of the *post-verify* strategy.

5.4 Case studies

Table 4: Comparison of mean average accepted tokens per second of vanilla SD methods and PEARL.

Methods	CodeLlama 7B&34B	CodeLlama 7B&70B	DeepSeek 1.3B&33B	DeepSeek 6.7B&33B
SD	5.27	8.32	7.23	5.69
PEARL	27.95	26.53	29.65	39.90

5.4.1 Optimal results of the window size γ

As mentioned in Section 4.1, PEARL eliminates the need for tuning the hyperparameter of window size and can theoretically predetermine the optimal value of γ. In this section, we conduct experiments with various gamma values in our PEARL to demonstrate the impact of the optimal γ on accelerating LLM inference. As shown in Table 5, we can observe that with the predetermined optimal γ, PEARL achieves the maximum inference acceleration. In certain cases involving suboptimal γ, the inference acceleration improvement in PEARL is less significant, underscoring the importance of predetermining the gamma value.

Table 5: Results on the optimal γ of different model combinations on HumanEval.

γ	CodeLlama 7&34 (c=3)	CodeLlama 7&70 (c=5)	DeepSeek 6.7&33 (c=3)
2	48.67	25.07	41.47
3	**65.59**	38.57	**65.12**
4	63.25	49.07	65.07
5	57.26	**58.10**	64.42
6	58.64	57.05	63.60

Table 6: Results on the optimal γ for different tasks of the same model combination. (c=5)

γ	HumanEval	GSM8K	MT-Bench
3	31.22	30.99	29.71
4	40.43	37.08	34.62
5	**46.03**	**42.20**	**37.96**
6	40.91	38.66	33.75
7	39.60	35.92	30.76

5.4.2 Mean accepted tokens

In Section 4.2, we theoretically demonstrate that the expected number of accepted tokens in PEARL exceeds that of the vanilla SD method. To further illustrate the real mean accepted tokens in PEARL under real-world complex conditions, we conduct experiments on the HumanEval, GSM8K, and MT-Bench datasets. As shown in Table 6, we still empirically observe that that PEARL obtains more accepted tokens compared to vanilla SD methods, which further demonstrates the effectiveness of the PEARL framework. Specifically, PEARL achieves the max number of mean accepted tokens to **39.9**, which significantly outperforms vanilla SD methods by a large margin. These results demonstrate that our PEARL can fully exploit the inference ability of the draft model for further acceleration.

6 Related Work

Transformer inference acceleration. There exists extensive works for transformer inference acceleration. This includes efforts of model compression [Zhu et al., 2023], efficient architecture design [Chitty-Venkata and Somani, 2022], and hardware optimization and implementation [Dao et al., 2022]. Model compression methods such as quantization [Choi et al., 2018], knowledge distillation [Hinton et al., 2015], and structure pruning [Han et al., 2015] aim at reducing the number of computational operations. Efficient architecture design is proposed to develop lightweight transformer architectures. Hardware optimization and implementation is proposed for efficient execution to fully exploit the hardware devices. These methods have achieved great success, while they are orthogonal to speculative decoding algorithms, which can be integrated for further speedup.

Draft-then-verify framework. While SD exhibits great acceleration effectiveness and lossless generalization quality, it remains a challenge to find a compact draft model with high distribution alignment. Some works focus on removing the necessity of the draft model. Self-speculative decoding [Zhang et al., 2023] proposes to skip some intermediate layers of the target model for drafting. Medusa [Cai et al., 2024] adds extra decoding heads at the top of the target model to generate drafts. Lookahead decoding[Fu et al., 2024] caches the generation trajectory (n-grams) as the drafts. DistillSpec [Zhou et al., 2023] utilizes distillation method to identify a compact draft model. Ouroboros [Zhao et al., 2024] combines the standard SD and lookahead decoding to generate more precise and longer drafts. Besides these works, SpecInfer [Miao et al., 2023] proposes tree attention, which is widely used to verify more drafts and increase the acceptance rate. However, all of them do not address the parallelism issue. From this perspective, our PEARL is orthogonal to these methods and can be integrated with these methods, which is left as a future work.

7 Conclusion and Future Work

Limitations. As our PEARL is a parallel acceleration framework, it remains a challenge to schedule the GPU resources to avoid resource competitions.

Conclusion. In this paper, we propose a novel inference acceleration framework, called PEARL, which significantly improves LLM inference efficiency. PEARL consists of two simple and effective strategies, i.e., *pre-verify* and *post-verify*, which effectively alleviates the mutual waiting problem with parallelism and adaptive draft length. Moreover, We theoretically derive the optimal window size and the mean accepted tokens of our PEARL, which demonstrates the effectiveness of our PEARL. Extensive experiments demonstrate that our proposed PEARL outperforms existing state-of-the-art methods on various text generation benchmarks.

Future work. For future research, we aim to integrate PEARL with existing accelerated inference methods to explore more efficient and resource-friendly acceleration approaches for LLM inference. Hopefully, PEARL will facilitate the future development of LLM inference acceleration.

References

Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.

The claude 3 model family: Opus, sonnet, haiku. URL `https://api.semanticscholar.org/CorpusID:268232499`.

Rishi Bommasani, Drew A Hudson, Ehsan Adeli, Russ Altman, Simran Arora, Sydney von Arx, Michael S Bernstein, Jeannette Bohg, Antoine Bosselut, Emma Brunskill, et al. On the opportunities and risks of foundation models. *arXiv preprint arXiv:2108.07258*, 2021.

Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. Llama: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*, 2023.

Diksha Khurana, Aditya Koli, Kiran Khatter, and Sukhdev Singh. Natural language processing: State of the art, current trends and challenges. *Multimedia tools and applications*, 82(3):3713–3744, 2023.

Tianle Cai, Yuhong Li, Zhengyang Geng, Hongwu Peng, Jason D Lee, Deming Chen, and Tri Dao. Medusa: Simple llm inference acceleration framework with multiple decoding heads. *arXiv preprint arXiv:2401.10774*, 2024.

Yichao Fu, Peter Bailis, Ion Stoica, and Hao Zhang. Break the sequential dependency of llm inference using lookahead decoding. *arXiv preprint arXiv:2402.02057*, 2024.

Jun Zhang, Jue Wang, Huan Li, Lidan Shou, Ke Chen, Gang Chen, and Sharad Mehrotra. Draft & verify: Lossless large language model acceleration via self-speculative decoding. *arXiv preprint arXiv:2309.08168*, 2023.

Yongchao Zhou, Kaifeng Lyu, Ankit Singh Rawat, Aditya Krishna Menon, Afshin Rostamizadeh, Sanjiv Kumar, Jean-François Kagy, and Rishabh Agarwal. Distillspec: Improving speculative decoding via knowledge distillation. *arXiv preprint arXiv:2310.08461*, 2023.

Weilin Zhao, Yuxiang Huang, Xu Han, Chaojun Xiao, Zhiyuan Liu, and Maosong Sun. Ouroboros: Speculative decoding with large model enhanced drafting. *arXiv preprint arXiv:2402.13720*, 2024.

Xupeng Miao, Gabriele Oliaro, Zhihao Zhang, Xinhao Cheng, Zeyu Wang, Rae Ying Yee Wong, Zhuoming Chen, Daiyaan Arfeen, Reyna Abhyankar, and Zhihao Jia. Specinfer: Accelerating generative llm serving with speculative inference and token tree verification. *arXiv preprint arXiv:2305.09781*, 2023.

Baptiste Roziere, Jonas Gehring, Fabian Gloeckle, Sten Sootla, Itai Gat, Xiaoqing Ellen Tan, Yossi Adi, Jingyu Liu, Tal Remez, Jérémy Rapin, et al. Code llama: Open foundation models for code. *arXiv preprint arXiv:2308.12950*, 2023.

Yaniv Leviathan, Matan Kalman, and Yossi Matias. Fast inference from transformers via speculative decoding. In *International Conference on Machine Learning*, pages 19274–19286. PMLR, 2023.

Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.

Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.

Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, et al. Judging llm-as-a-judge with mt-bench and chatbot arena. *Advances in Neural Information Processing Systems*, 36, 2024.

Daya Guo, Qihao Zhu, Dejian Yang, Zhenda Xie, Kai Dong, Wentao Zhang, Guanting Chen, Xiao Bi, Y Wu, YK Li, et al. Deepseek-coder: When the large language model meets programming–the rise of code intelligence. *arXiv preprint arXiv:2401.14196*, 2024.

Charlie Chen, Sebastian Borgeaud, Geoffrey Irving, Jean-Baptiste Lespiau, Laurent Sifre, and John Jumper. Accelerating large language model decoding with speculative sampling. *arXiv preprint arXiv:2302.01318*, 2023.

Xunyu Zhu, Jian Li, Yong Liu, Can Ma, and Weiping Wang. A survey on model compression for large language models. *arXiv preprint arXiv:2308.07633*, 2023.

Krishna Teja Chitty-Venkata and Arun K Somani. Neural architecture search survey: A hardware perspective. *ACM Computing Surveys*, 55(4):1–36, 2022.

Tri Dao, Dan Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. Flashattention: Fast and memory-efficient exact attention with io-awareness. *Advances in Neural Information Processing Systems*, 35:16344–16359, 2022.

Jungwook Choi, Zhuo Wang, Swagath Venkataramani, Pierce I-Jen Chuang, Vijayalakshmi Srinivasan, and Kailash Gopalakrishnan. Pact: Parameterized clipping activation for quantized neural networks. *arXiv preprint arXiv:1805.06085*, 2018.

Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015.

Song Han, Huizi Mao, and William J Dally. Deep compression: Compressing deep neural networks with pruning, trained quantization and huffman coding. *arXiv preprint arXiv:1510.00149*, 2015.

A Algorithm of PEARL

Here, we give the whole algorithm of our PEARL in details in Algorithm. 2.

Algorithm 2 Parallel Speculative Decoding with Adaptive Draft Length.

Input: the draft model M_q, the target model M_p, the input prefix $\mathbf{x}$, the max generate tokens L, the window size γ.

    ▷ The *pre-verify* strategy is used first.
1: Initialization: mode ← "pre-verify"
2: **while** $len(\mathbf{x}) < L$ **do**
3:    **if** mode = "pre-verify" **then**
4:       ▷ Pre-verify strategy
5:       **for** $i = 1$ to γ **do**
6:          $q_i \leftarrow M_q(\mathbf{x} + [x_1, ..., x_{i-1}])$
7:          $x_i \sim q_i$
8:       **end for**
9:       ▷ running the target model in parallel to verify the first draft token in advance.
10:       $p \leftarrow M_p(\mathbf{x})$
11:       **if** $r \sim U(0,1) \leq \frac{p[x_1]}{q_1[x_1]}$ **then**
12:          ✓ accept the first token
13:          $\mathbf{x} \leftarrow \mathbf{x} + [x_1, ..., x_\gamma]$
14:          mode ← "post-verify"
15:       **else**
16:          × reject the first token
17:          $y \sim norm(max(0, p - q_1))$
18:          $\mathbf{x} \leftarrow \mathbf{x} + [y]$
19:          mode ← "pre-verify"
20:       **end if**
21:    **else**
22:       ▷ Post-verify strategy
23:       $\mathbf{x}, [x_1, x_2, ..., x_\gamma] \leftarrow \mathbf{x}$ ▷ split the prefix to get the last γ draft tokens
24:       **for** $i = \gamma + 1$ to 2γ **do**
25:          ▷ running the draft model in parallel to continue drafting.
26:          $q_i \leftarrow M_q(\mathbf{x} + [x_1, ..., x_{i-1}])$
27:          $x_i \sim q_i$
28:       **end for**
29:       $p_1, p_2, ..., p_\gamma \leftarrow M_p(\mathbf{x} + [x_1]), M_p(\mathbf{x} + [x_1, x_2]), ..., M_p(\mathbf{x} + [x_1, ..., x_\gamma])$
30:       retrival $q_1, q_2, ..., q_\gamma$ from the cache
31:       $r_1 \sim U(0,1), ..., r_\gamma \sim U(0,1)$
32:       $n \leftarrow \min(\{i - 1 | 1 \leq i \leq \gamma, r_i > \frac{p_i[x_i]}{q_i[x_i]}\} \cup \{\gamma\})$
33:       **if** $n = \gamma$ **then**
34:          ✓ accept all draft tokens
35:          $\mathbf{x} \leftarrow \mathbf{x} + [x_1, ..., x_{2\gamma}]$
36:          mode ← "post-verify"
37:       **else**
38:          × reject someone
39:          $y \sim norm(max(0, p_{n+1} - q_{n+1}))$
40:          $\mathbf{x} \leftarrow \mathbf{x} + [x_1, ..., x_n, y]$
41:          mode ← "pre-verify"
42:       **end if**
43:    **end if**
44: **end while**

B Proof of Theorem 1 and Theorem 2

As illustrated in Section 4, we have two theorems to demonstrate the effectiveness of our PEARL. Here we give the proof of these two theorems.

Theorem 1 *Given a draft model M_q and a target model M_p, the optimal value of the window size γ is the ratio of the running speed of the draft model and the target model, i.e.,*

$$\gamma' = \arg\max_{\gamma} \ PEARL(\gamma) = c. \tag{5}$$

Proof. We discuss the situation that $\gamma < c$ and $\gamma > c$.

If γ is smaller than c, the draft model will wait for the target model until it finishes the verification. The draft model could have generated more draft tokens to be verified without extra time consumed, That is to say, the situation $\gamma < c$ decreases the number of verified tokens, where PEARL needs to take more steps for inference.

If γ is larger than c, the target model will wait for the draft model until it finishes generating draft tokens. Theoretically, the drafting phase takes $\max(\gamma t, ct)$ and the verification phase takes ct, while the verification phase can verify γ tokens. Therefore, the verification speed is $\frac{\max(\gamma t, ct) + ct}{\gamma}$. As $\gamma > c$, the verification speed is $\frac{(\gamma+c)}{\gamma}t$. Obviously, the optimal value of γ is c.

□

Theorem 2 *Assuming the acceptance rate of each draft token is α, and α is i.i.d., the expectation of the number of accepted tokens of PEARL is*

$$E(\#accepted\ tokens) = \frac{1}{1-\alpha} + 1. \tag{6}$$

Proof. While vanilla speculative decoding can decode $\gamma + 1$ tokens at most, our PEARL can decode infinity tokens due to the adaptive draft length, therefore the expectation of accepted tokens is given by:

$$\begin{aligned} E &= \sum_{k=0}^{\infty} kP(acc = k) \\ &= (1-\alpha)\sum_{k=0}^{\infty} k\alpha^k \end{aligned} \tag{7}$$

Let $S = \sum_{k=0}^{\infty} k\alpha^k$, $\alpha S = \sum_{k=1}^{\infty}(k-1)\alpha^k$, we have:

$$\begin{aligned} E &= (1-\alpha)S = S - \alpha S \\ &= \sum_{k=0}^{\infty} k\alpha^k - \sum_{k=1}^{\infty}(k-1)\alpha^k \\ &= \frac{1}{1-\alpha} \end{aligned} \tag{8}$$

Counting the additional token generated by the target model, the expectation is:

$$E = \frac{1}{1-\alpha} + 1 \tag{9}$$

□

C Evaluation Details

C.1 Model Configurations

We select some representative models for evaluation, including Llama 2 Touvron et al. [2023], Codellama Roziere et al. [2023] and Deepseek-Coder Guo et al. [2024]. We summarize the model configuration in Table 7. In our experiments, all models are loaded in the precision of bfloat-16. Our PEARL does not introduce any additional training, and directly uses these models to evaluate our algorithm. The running speed is measured on the code generation tasks.

Table 7: Detailed model configurations.

Models	Layers	dim	FFN dim	speed (tok/s)
Codellama-7B	32	4096	11008	73.38
Codellama-34B	48	8192	22016	27.96
Codellama-70B	80	8192	28672	15.33
Deepseek-1.3B	24	2048	5504	102.36
Deepseek-6.7B	32	4096	11008	73.17
Deepseek-33B	62	7168	19200	23.64
Llama-2-7B	32	4096	11008	71.10
Llama-2-70B	80	8192	28672	15.30

C.2 Evaluation Details

All of our experiments including latency measurement, ablation studies, and case studies are conducted on NVIDIA H800 GPUs. For models with size of 1.3B and 7B, we put them on a single H800, while 34B models are deployed on 2 H800, 70B models are deployed on 3 H800. For inference, we use batch size 1, which is commonly used in other speculative decoding works. For the compared baselines, including Lookahead decoding and Ouroboros, we reproduce the results of them on the code generation tasks with the default parameters as described in their paper or code. When evaluating these methods, the model configuration and GPU usage is the same as our PEARL. For HumanEval, we set the number of max generated tokens as 1024. For the other two tasks, we set the number of max generated tokens as 512.

As our PEARL is a parallel inference acceleration framework, we implement the parallel algorithm in accelerate, which can be further optimized with other parallel techniques. We leave this as a potential future work to acquire more acceleration.